

*Raphaël Couturier*

---

# ***Designing Scientific Applications on GPUs***



---

## ***Foreward***

I am delighted to introduce the first book on Multimedia Data Mining. When I came to know about this book project undertaken by two of the most active young researchers in the field, I was pleased that this book is coming in early stage of a field that will need it more than most fields do. In most emerging research fields, a book can play a significant role in bringing some maturity to the field. Research fields advance through research papers. In research papers, however, only a limited perspective could be provided about the field, its application potential, and the techniques required and already developed in the field. A book gives such a chance. I liked the idea that there will be a book that will try to unify the field by bringing in disparate topics already available in several papers that are not easy to find and understand. I was supportive of this book project even before I had seen any material on it. The project was a brilliant and a bold idea by two active researchers. Now that I have it on my screen, it appears to be even a better idea.

Multimedia started gaining recognition in 1990s as a field. Processing, storage, communication, and capture and display technologies had advanced enough that researchers and technologists started building approaches to combine information in multiple types of signals such as audio, images, video, and text. Multimedia computing and communication techniques recognize correlated information in multiple sources as well as insufficiency of information in any individual source. By properly selecting sources to provide complementary information, such systems aspire, much like human perception system, to create a holistic picture of a situation using only partial information from separate sources.

Data mining is a direct outgrowth of progress in data storage and processing speeds. When it became possible to store large volume of data and run different statistical computations to explore all possible and even unlikely correlations among data, the field of data mining was born. Data mining allowed people to hypothesize relationships among data entities and explore support for those. This field has been put to applications in many diverse domains and keeps getting more applications. In fact many new fields are direct outgrowth of data mining and it is likely to become a powerful computational tool.



---

## *Preface*

Approximately 17 million people in the USA (6% of the population) and 140 million people worldwide (this number is expected to rise to almost 300 million by the year 2025) suffer from *diabetes mellitus*. Currently, there are a few dozens of commercialised devices for detecting blood glucose levels [1]. However, most of them are invasive. The development of a noninvasive method would considerably improve the quality of life for diabetic patients, facilitate their compliance for glucose monitoring, and reduce complications and mortality associated with this disease. Noninvasive and continuous monitoring of glucose concentration in blood and tissues is one of the most challenging and exciting applications of optics in medicine. The major difficulty in development and clinical application of optical noninvasive blood glucose sensors is associated with very low signal produced by glucose molecules. This results in low sensitivity and specificity of glucose monitoring by optical methods and needs a lot of efforts to overcome this difficulty.

A wide range of optical technologies have been designed in attempts to develop robust noninvasive methods for glucose sensing. The methods include infrared absorption, near-infrared scattering, Raman, fluorescent, and thermal gradient spectroscopies, as well as polarimetric, polarization heterodyning, photonic crystal, optoacoustic, optothermal, and optical coherence tomography (OCT) techniques [1-31].

For example, the polarimetric quantification of glucose is based on the phenomenon of optical rotatory dispersion, whereby a chiral molecule in an aqueous solution rotates the plane of linearly polarized light passing through the solution. The angle of rotation depends linearly on the concentration of the chiral species, the pathlength through the sample, and the molecule specific rotation. However, polarization sensitive optical technique makes it difficult to measure *in vivo* glucose concentration in blood through the skin because of the strong light scattering which causes light depolarization. For this reason, the anterior chamber of the eye has been suggested as a sight well suited for polarimetric measurements, since scattering in the eye is generally very low compared to that in other tissues, and a high correlation exists between the glucose in the blood and in the aqueous humor. The high accuracy of anterior eye chamber measurements is also due to the low concentration of optically active aqueous proteins within the aqueous humor.

On the other hand, the concept of noninvasive blood glucose sensing using the scattering properties of blood and tissues as an alternative to spectral

absorption and polarization methods for monitoring of physiological glucose concentrations in diabetic patients has been under intensive discussion for the last decade. Many of the considered effects, such as changing of the size, refractive index, packing, and aggregation of RBC under glucose variation, are important for glucose monitoring in diabetic patients. Indeed, at physiological concentrations of glucose, ranging from 40 to 400 mg/dl, the role of some of the effects may be modified, and some other effects, such as glucose penetration inside the RBC and the followed hemoglobin glycation, may be important [30-32].

Noninvasive determination of glucose was attempted using light scattering of skin tissue components measured by a spatially-resolved diffuse reflectance or NIR frequency-domain reflectance techniques. Both approaches are based on change in glucose concentration, which affects the refractive index mismatch between the interstitial fluid and tissue fibers, and hence reduces scattering coefficient. A glucose clamp experiment showed that reduced scattering coefficient measured in the visible range qualitatively tracked changes in blood glucose concentration for the volunteer with diabetes studied.

---

## *List of Figures*

1.1	Comparison of number of cores in a CPU and in a GPU. . . .	6
1.2	Comparison of low latency of CPU and high throughput of GPU. . . . .	6
1.3	Scalability of GPU. . . . .	8
1.4	Memory hierarchy of a GPU. . . . .	10





---

## *List of Tables*



---

# Contents

<b>I</b>	<b>This is a Part</b>	<b>1</b>
<b>1</b>	<b>Presentation of the GPU architecture and of the Cuda environment</b>	<b>3</b>
	<i>Raphaël Couturier</i>	
1.1	Introduction . . . . .	3
1.2	Brief history of Video Card . . . . .	3
1.3	GPGPU . . . . .	4
1.4	Architecture of current GPUs . . . . .	5
1.5	Kinds of parallelism . . . . .	7
1.6	Cuda Multithreading . . . . .	7
1.7	Memory hierarchy . . . . .	9
1.8	Conclusion . . . . .	10
	Bibliography . . . . .	11
<b>2</b>	<b>Introduction to Cuda</b>	<b>13</b>
	<i>Raphaël Couturier</i>	
2.1	Introduction . . . . .	13
2.2	First example . . . . .	13
2.3	Second example: using CUBLAS . . . . .	16
2.4	Third example: matrix-matrix multiplication . . . . .	19
2.5	Conclusion . . . . .	22
	Bibliography . . . . .	22





## Part I

This is a Part



# Chapter 1

---

## *Presentation of the GPU architecture and of the Cuda environment*

**Raphaël Couturier**

*Femto-ST Institute, University of Franche-Comte*

1.1	Introduction .....	3
1.2	Brief history of Video Card .....	3
1.3	GPGPU .....	4
1.4	Architecture of current GPUs .....	5
1.5	Kinds of parallelism .....	7
1.6	Cuda Multithreading .....	7
1.7	Memory hierarchy .....	9
1.8	Conclusion .....	10
	Bibliography .....	10

---

### **1.1 Introduction**

This chapter introduces the Graphics Processing Unit (GPU) architecture and all the concepts needed to understand how GPUs work and can be used to speed up the execution of some algorithms. First of all this chapter gives a brief history of the development of Graphics card until they have been used in order to make general purpose computation. Then the architecture of a GPU is illustrated. There are many fundamental differences between a GPU and a tradition processor. In order to benefit from the power of a GPU, a Cuda programmer needs to use threads. They have some particularities which enable the Cuda model to be efficient and scalable when some constraints are addressed.

---

### **1.2 Brief history of Video Card**

Video cards or Graphics cards have been introduced in personal computers to produce high quality graphics faster than classical Central Processing Units (CPU) and to alleviate CPU from this task. In general, display tasks are

very repetitive and very specific. Hence, some manufacturers have produced more and more sophisticated video cards, providing 2D accelerations then 3D accelerations, then some light transforms. Video cards own their own memory to perform their computation. For at least two decades, every personal computer has had a video card which is simple for desktop computers or which provides many accelerations for game and/or graphic oriented computers. In the latter case, graphic cards may be more expensive than a CPU.

Since 2000, video cards have allowed users to apply arithmetic operations simultaneously on a sequence of pixels, also later called stream processing. In this case, the information of the pixels (color, location and other information) are combined in order to produce a pixel color that can be displayed on a screen. Simultaneous computations are provided by shaders which calculate rendering effects on graphics hardware with a high degree of flexibility. These shaders handles the stream data with pipelines.

Some researchers tried to apply those operations on other data, representing something different from pixels, and consequently this resulted in the first uses of video cards for performing general purpose computation. The programming model was not easy to use at all and was very dependent of the hardware constraints. More precisely it consisted in using either DirectX or OpenGL functions providing an interface to some classical operations for videos operations (memory transfers, texture manipulation, ...). Floating point operations were most of the time unimaginable. Obviously when something went wrong, programmers had no way (and neither the tools) to detect it.

---

### 1.3 GPGPU

In order to benefit from the computing power of more recent video cards, Cuda was first proposed in 2007 by NVidia. It unifies the programming model for some of their most performant video cards. Cuda [3] has quickly been considered by the scientific community as a great advance for general purpose graphics processing unit (GPGPU) computing. Of course other programming models have been proposed. The other well-known alternative is OpenCL which aims at proposing an alternative to Cuda and which is multi-platform and portable. This is a great advantage since it is even possible to execute OpenCL programs on traditional CPUs. The main drawback is that it is less tight with the hardware and consequently sometimes provides less efficient programs. Moreover, Cuda benefits from more mature compilation and optimization procedures. Other less known environments have been proposed, but most of them have been stopped, for example we can cite: FireStream by ATI which is not maintained anymore and replaced by OpenCL, BrookGPU by Stanford University [1]. Another environment based on pragma (insertion of pragma directives inside the code to help the compiler to generate efficient



code) is called OpenACC. For a comparison with OpenCL, interested readers may refer to [2].

---

## 1.4 Architecture of current GPUs

The architecture of current GPUs is constantly evolving. Nevertheless some trends remain constant throughout this evolution. Processing units composing a GPU are far more simple than a traditional CPU but it is much easier to integrate many computing units inside a GPU card than to do so with many cores inside a CPU. This is due to the fact that the cores of a GPU are simpler than the cores of a CPU. In 2012, the most powerful GPUs own more than 500 cores and the most powerful CPUs have 8 cores. Figure 1.1 shows the number of cores inside a CPU and inside a GPU. In fact, in a current NVidia GPU, there are multiprocessors which have 32 cores (for example on Fermi cards). The core clock of CPU is generally around 3GHz and the one of GPU is about 1.5GHz. Although the core clock of GPU cores is slower, the amount of cores inside a GPU provides more computational power. This measure is commonly represented by the number of floating point operation per seconds. Nowadays the most powerful GPUs provide more than 1TFlops, i.e.  $10^{12}$  floating point operations per second. Nevertheless GPUs are very efficient to perform some operations but not all kinds of operations. They are very efficient to execute repetitive work in which only the data change. It is important to keep in mind that multiprocessors inside a GPU have 32 cores. Later we will see that these 32 cores need to do the same work to get maximum performance.

On the most powerful GPU cards, called Fermi, multiprocessors are called streaming multiprocessors (SM). Each SM contains 32 cores and is able to perform 32 floating points or integer operations on 32 bits numbers per clock or 16 floating points on 64 bits number per clock. SMs have their own registers, execution pipelines and caches. On Fermi architecture, there are 64Kb shared memory + L1 cache and 32,536 32bits registers per SM. More precisely the programmer can decide what amount of shared memory and L1 cache SM can use. The constraint is that the sum of both amounts should be less or equal to 64Kb.

Threads are used to benefit from the important number of cores of a GPU. Those threads are different from traditional threads for CPU. In chapter 2, some examples of GPU programming will explicit the details of the GPU threads. However, threads are gathered into blocks of 32 threads, called “warps”. Those warps are important when designing an algorithm for GPU.

Another big difference between CPU and GPU is the latency of memory. In CPU, everything is optimized to obtain a low latency architecture. This is possible through the use of cache memories. Moreover, nowadays CPUs perform many performance optimizations such as speculative execution which

roughly speaking consists in executing a small part of code in advance even if later this work reveals itself to be useless. On the contrary, GPUs do not have low latency memory. In comparison GPUs have small cache memories. Nevertheless the architecture of GPUs is optimized for throughput computation and it takes into account the memory latency.

Figure 1.2 illustrates the main difference of memory latency between a CPU and a GPU. In a CPU, tasks “ $t_i$ ” are executed one by one with a short

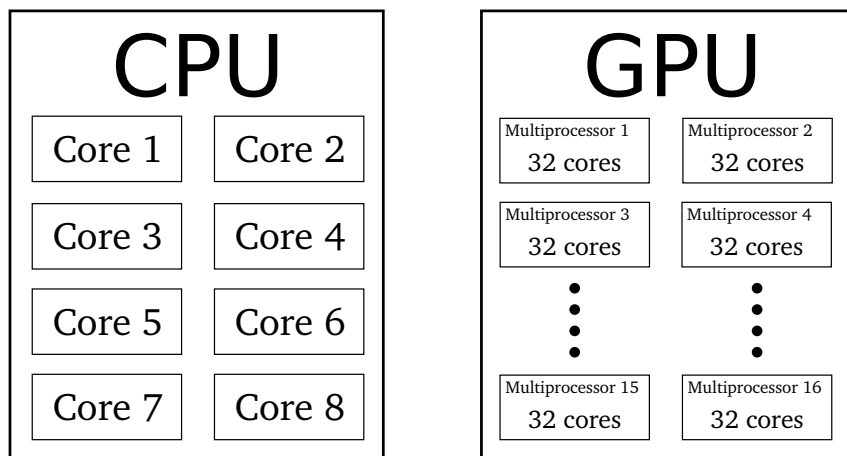



FIGURE 1.1: Comparison of number of cores in a CPU and in a GPU.

CPU: optimized for low latency



GPU: optimized for high throughput




FIGURE 1.2: Comparison of low latency of CPU and high throughput of GPU.

memory latency to get the data to process. After some tasks, there is a context switch that allows the CPU to run concurrent applications and/or multi-threaded applications. Memory latencies are longer in a GPU, the the principle to obtain a high throughput is to have many tasks to compute. Later we will see that those tasks are called threads with Cuda. With this principle, as soon as a task is finished the next one is ready to be executed while the wait for data for the previous task is overlapped by computation of other tasks.

---

## 1.5 Kinds of parallelism

Many kinds of parallelism are amiable according to the type of hardware. Roughly speaking, there are three classes of parallelism: instruction-level parallelism, data parallelism and task parallelism.

Instruction-level parallelism consists in re-ordering some instructions in order to execute some of them in parallel without changing the result of the code. In modern CPUs, instruction pipelines allow processor to execute instructions faster. With a pipeline a processor can execute multiple instructions simultaneously due to the fact that the output of a task is the input of the next one.

Data parallelism consists in executing the same program with different data on different computing units. Of course, no dependency should exist between the data. For example, it is easy to parallelize loops without dependency using the data parallelism paradigm. This paradigm is linked with the Single Instructions Multiple Data (SIMD) architecture. This is the kind of parallelism provided by GPUs.

Task parallelism is the common parallelism achieved out on clusters and grids and high performance architectures where different tasks are executed by different computing units.

---

## 1.6 Cuda Multithreading

The data parallelism of Cuda is more precisely based on the Single Instruction Multiple Thread (SIMT) model. This is due to the fact that a programmer accesses to the cores by the intermediate of threads. In the Cuda model, all cores execute the same set of instructions but with different data. This model has similarities with the vector programming model proposed for vector machines through the 1970s into the 90s, notably the various Cray platforms. On the Cuda architecture, the performance is led by the use of a huge number of

threads (from thousands up to millions). The particularity of the model is that there is no context switching as in CPUs and each thread has its own registers. In practice, threads are executed by SM and are gathered into groups of 32 threads. Those groups are called “warps”. Each SM alternatively executes “active warps” and warps becoming temporarily inactive due to waiting of data (as shown in Figure 1.2).

The key to scalability in the Cuda model is the use of a huge number of threads. In practice, threads are not only gathered in warps but also in thread blocks. A thread block is executed by only one SM and it cannot migrate. The typical size of a thread block is a number power of two (for example: 64, 128, 256 or 512).

In this case, without changing anything inside a Cuda code, it is possible to run your code with a small Cuda device or the most performing Tesla Cuda cards. Blocks are executed in any order depending on the number of SMs available. So the programmer must conceive its code having this issue in mind. This independence between thread blocks provides the scalability of Cuda codes.

A kernel is a function which contains a block of instructions that are ex-

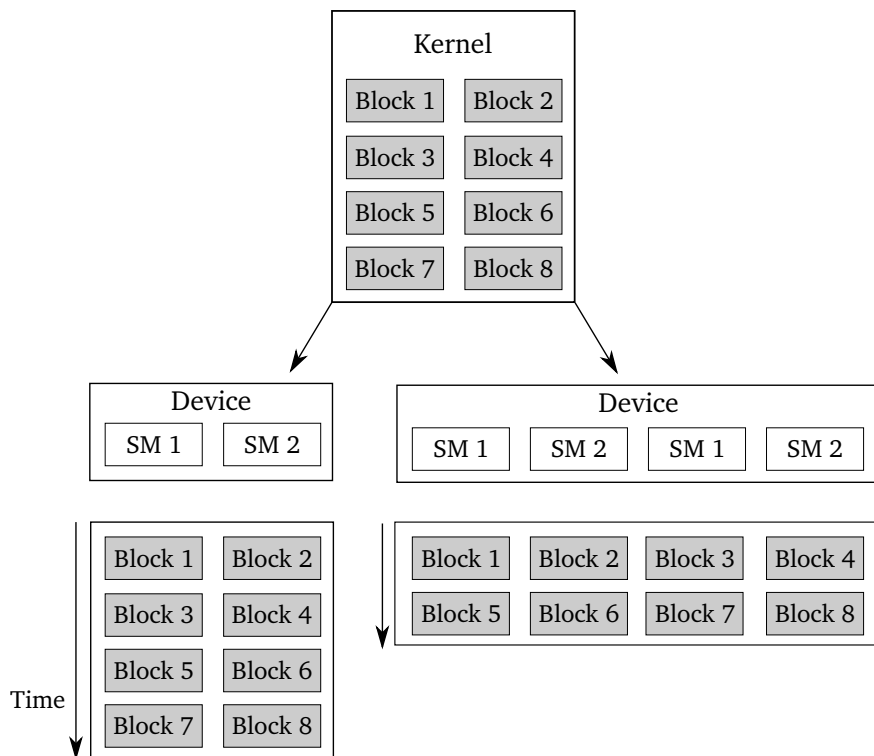


FIGURE 1.3: Scalability of GPU.

executed by the threads of a GPU. When the problem considered is a two dimensional or three dimensional problem, it is possible to group thread blocks into a grid. In practice, the number of thread blocks and the size of thread blocks is given as parameters to each kernel. Figure 1.3 illustrates an example of a kernel composed of 8 thread blocks. Then this kernel is executed on a small device containing only 2 SMs. So in this case, blocks are executed 2 by 2 in any order. If the kernel is executed on a larger Cuda device containing 4 SMs, blocks are executed 4 by 4 simultaneously. The execution times should be approximately twice faster in the latter case. Of course, that depends on other parameters that will be described later.

Thread blocks provide a way to cooperation in the sense that threads of the same block cooperatively load and store blocks of memory they all use. Synchronizations of threads in the same block are possible (but not between threads of different blocks). Threads of the same block can also share results in order to compute a single result. In chapter 2, some examples will explicit that.

---

## 1.7 Memory hierarchy

The memory hierarchy of GPUs is different from the CPUs one. In practice, there are registers, local memory, shared memory, cache memory and global memory.

As previously mentioned each thread can access its own registers. It is important to keep in mind that the number of registers per block is limited. On recent cards, this number is limited to 64Kb per SM. Access to registers is very fast, so it is a good idea to use them whenever possible.

Likewise each thread can access local memory which, in practice, is much slower than registers. Local memory is automatically used by the compiler when all the registers are occupied. So the best idea is to optimize the use of registers even if this implies to reduce the number of threads per block.

Shared memory allows cooperation between threads of the same block. This kind of memory is fast because it requires to be manipulated manually and its size is limited. It is accessible during the execution of a kernel. So the principle is to fill the shared memory at the start of the kernel with global data that are used very frequently, then threads can access it for their computation. They can obviously change the content of this shared memory either with computation or load of other data and they can store its content in the global memory. So shared memory can be seen as a cache memory manageable manually. This requires obviously an effort from the programmer.

On recent cards, the programmer may decide what amount of cache memory and shared memory is attributed to a kernel. The cache memory is a L1 cache which is directly managed by the GPU. Sometimes, this cache pro-

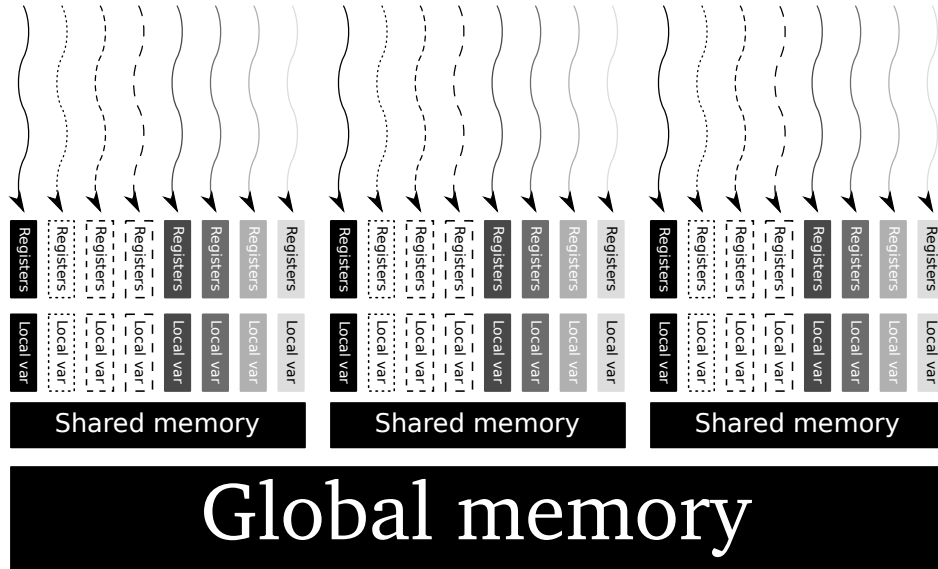


FIGURE 1.4: Memory hierarchy of a GPU.

vides very efficient result and sometimes the use of shared memory is a better solution.

Figure 1.4 illustrates the memory hierarchy of a GPU. Threads are represented on the top of the figure. They can access to their own registers and their local memory. Threads of the same block can access to the shared memory of this block. The cache memory is not represented here but it is local to a thread. Then each block can access to the global memory of the GPU.

## 1.8 Conclusion

In this chapter, a brief presentation of the video card, which has later been used to perform computation, has been given. The architecture of a GPU has been illustrated focusing on the particularity of GPUs in term of parallelism, memory latency and threads. In order to design an efficient algorithm for GPU, it is essential to have all these parameters in mind.

## Bibliography

- [1] I. Buck, T. Foley, D. Horn, J. Sugerman, K. Fatahalian, M. Houston, and P. Hanrahan. Brook for GPUs: stream computing on graphics hardware. *ACM Transactions on Graphics*, 23(3):777–786, August 2004.
- [2] B. Cloutier, B. K. Muite, and P. Rigge. A comparison of cpu and gpu performance for fourier pseudospectral simulations of the navier-stokes, cubic nonlinear schrödinger and sine gordon equations. *Computational Physics (physics.comp-ph)*, 2012.
- [3] NVIDIA Corporation. NVIDIA CUDA C programming guide, 2011. Version 4.0.





# Chapter 2

---

## Introduction to Cuda

**Raphaël Couturier**

*Femto-ST Institute, University of Franche-Comte*

2.1	Introduction .....	13
2.2	First example .....	13
2.3	Second example: using CUBLAS .....	16
2.4	Third example: matrix-matrix multiplication .....	19
2.5	Conclusion .....	22
	Bibliography .....	22

---

### 2.1 Introduction

In this chapter we give some simple examples of Cuda programming. The goal is not to provide an exhaustive presentation of all the functionalities of Cuda but rather to give some basic elements. Of course, readers that do not know Cuda are invited to read other books that are specialized on Cuda programming (for example: [2]).

---

### 2.2 First example

This first example is intended to show how to build a very simple program with Cuda. Its goal is to perform the sum of two arrays and put the result into a third array. A Cuda program consists in a C code which calls Cuda kernels that are executed on a GPU. The listing of this code is in Listing 2.1.

As GPUs have their own memory, the first step consists in allocating memory on the GPU. A call to `cudaMalloc` allows to allocate memory on the GPU. The first parameter of this function is a pointer on a memory on the device (i.e. the GPU). In this example, `d_` is added on each variable allocated on the GPU, meaning this variable is on the GPU. The second parameter represents the size of the allocated variables, this size is expressed in bits.

In this example, we want to compare the execution time of the additions of two arrays in CPU and GPU. So for both these operations, a timer is created

to measure the time. Cuda proposes to manipulate timers quite easily. The first step is to create the timer, then to start it and at the end to stop it. For each of these operations a dedicated function is used.

In order to compute the same sum with a GPU, the first step consists in transferring the data from the CPU (considered as the host with Cuda) to the GPU (considered as the device with Cuda). A call to `cudaMemcpy` allows to copy the content of an array allocated in the host to the device when the fourth parameter is set to `cudaMemcpyHostToDevice`. The first parameter of the function is the destination array, the second is the source array and the third is the number of elements to copy (expressed in bytes).

Now the GPU contains the data needed to perform the addition. In sequential programming, such addition is achieved out with a loop on all the elements. With a GPU, it is possible to perform the addition of all the elements of the two arrays in parallel (if the number of blocks and threads per blocks is sufficient). In Listing 2.1 at the beginning, a simple kernel, called `addition` is defined to compute in parallel the summation of the two arrays. With Cuda, a kernel starts with the keyword `__global__` which indicates that this kernel can be called from the C code. The first instruction in this kernel is used to compute the variable `tid` which represents the thread index. This thread index is computed according to the values of the block index (called `blockIdx` in Cuda) and of the thread index (called `threadIdx` in Cuda). Blocks of threads and thread indexes can be decomposed into 1 dimension, 2 dimensions or 3 dimensions. According to the dimension of manipulated data, the appropriate dimension can be useful. In our example, only one dimension is used. Then using notation `.x` we can access to the first dimension (`.y` and `.z` respectively allow to access to the second and third dimension). The variable `blockDim` gives the size of each block.

**Listing 2.1: A simple example**

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <math.h>
#include <assert.h>
#include "cutil_inline.h"

const int nbThreadsPerBloc=256;

__global__
void addition(int size, int *d_C, int *d_A, int *d_B) {
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    if(tid<size) {
        d_C[tid]=d_A[tid]+d_B[tid];
    }
}

int main( int argc, char** argv)
{
```

```

if(argc!=2) {
    printf("usage: ex1 nb_components\n");
    exit(0);
}

int size=atoi(argv[1]);
int i;
int *h_arrayA=(int*)malloc(size*sizeof(int));
int *h_arrayB=(int*)malloc(size*sizeof(int));
int *h_arrayC=(int*)malloc(size*sizeof(int));
int *h_arrayCgpu=(int*)malloc(size*sizeof(int));
int *d_arrayA, *d_arrayB, *d_arrayC;

cudaMalloc((void**)&d_arrayA, size*sizeof(int));
cudaMalloc((void**)&d_arrayB, size*sizeof(int));
cudaMalloc((void**)&d_arrayC, size*sizeof(int));

for(i=0;i<size;i++) {
    h_arrayA[i]=i;
    h_arrayB[i]=2*i;
}

unsigned int timer_cpu = 0;
cutilCheckError(cutCreateTimer(&timer_cpu));
cutilCheckError(cutStartTimer(timer_cpu));
for(i=0;i<size;i++) {
    h_arrayC[i]=h_arrayA[i]+h_arrayB[i];
}
cutilCheckError(cutStopTimer(timer_cpu));
printf("CPU processing time : %f (ms) \n", cutGetTimerValue(
    timer_cpu));
cutDeleteTimer(timer_cpu);

unsigned int timer_gpu = 0;
cutilCheckError(cutCreateTimer(&timer_gpu));
cutilCheckError(cutStartTimer(timer_gpu));
cudaMemcpy(d_arrayA, h_arrayA, size * sizeof(int),
    cudaMemcpyHostToDevice);
cudaMemcpy(d_arrayB, h_arrayB, size * sizeof(int),
    cudaMemcpyHostToDevice);

int nbBlocs=(size+nbThreadsPerBloc-1)/nbThreadsPerBloc;
addition<<<nbBlocs,nbThreadsPerBloc>>>(size,d_arrayC,d_arrayA,
    d_arrayB);
cudaMemcpy(h_arrayCgpu,d_arrayC, size * sizeof(int),
    cudaMemcpyDeviceToHost);

cutilCheckError(cutStopTimer(timer_gpu));
printf("GPU processing time : %f (ms) \n", cutGetTimerValue(
    timer_gpu));
cutDeleteTimer(timer_gpu);

for(i=0;i<size;i++)
    assert(h_arrayC[i]==h_arrayCgpu[i]);

cudaFree(d_arrayA);
cudaFree(d_arrayB);

```

```
cudaFree(d_arrayC);  
free(h_arrayA);  
free(h_arrayB);  
free(h_arrayC);  
return 0;  
}
```

---

## 2.3 Second example: using CUBLAS

The Basic Linear Algebra Subprograms (BLAS) allows programmers to use efficient routines that are often required. Those routines are heavily used in many scientific applications and are optimized for vector operations, matrix-vector operations and matrix-matrix operations [1]. Some of those operations seem to be easy to implement with Cuda. Nevertheless, as soon as a reduction is needed, implementing an efficient reduction routine with Cuda is far from being simple. Roughly speaking, a reduction operation is an operation which combines all the elements of an array and extracts a number computed with all the elements. For example, a sum, a maximum or a dot product are reduction operations.

In this second example, we consider that we have two vectors  $A$  and  $B$ . First of all, we want to compute the sum of both vectors in a vector  $C$ . Then we want to compute the scalar product between  $1/C$  and  $1/A$ . This is just an example which has no direct interest except to show how to program it with Cuda.

Listing 2.2 shows this example with Cuda. The first kernel for the addition of two arrays is exactly the same as the one described in the previous example.

The kernel to compute the opposite of the elements of an array is very simple. For each thread index, the inverse of the array replaces the initial array.

In the main function, the beginning is very similar to the one in the previous example. First, the user is asked to define the number of elements. Then a call to `cublasCreate` allows to initialize the cublas library. It creates a handle. Then all the arrays are allocated in the host and the device, as in the previous example. Both arrays  $A$  and  $B$  are initialized. The CPU computation is performed and the time for this CPU computation is measured. In order to compute the same result on the GPU, first of all, data from the CPU need to be copied into the memory of the GPU. For that, it is possible to use cublas function `cublasSetVector`. This function has several arguments. More precisely, the first argument represents the number of elements to transfer, the second arguments is the size of each element, the third element represents the source of the array to transfer (in the GPU), the fourth is an offset between each element of the source (usually this value is set to

1), the fifth is the destination (in the GPU) and the last is an offset between each element of the destination. Then we call the kernel addition which computes the sum of all elements of arrays  $A$  and  $B$ . The inverse kernel is called twice, once to inverse elements of array  $C$  and once for  $A$ . Finally, we call the function `cublasDdot` which computes the dot product of two vectors. To use this routine, we must specify the handle initialized by Cuda, the number of elements to consider, then each vector is followed by the offset between every element. After the GPU computation, it is possible to check that both computation produce the same result.

Listing 2.2: A simple example with cublas

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <math.h>
#include <assert.h>
#include "cutil_inline.h"
#include <cublas_v2.h>

const int nbThreadsPerBloc=256;

__global__
void addition(int size, double *d_C, double *d_A, double *d_B) {
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    if(tid<size) {
        d_C[tid]=d_A[tid]+d_B[tid];
    }
}

__global__
void inverse(int size, double *d_x) {
    int tid = blockIdx.x * blockDim.x + threadIdx.x;
    if(tid<size) {
        d_x[tid]=1./d_x[tid];
    }
}

int main( int argc, char** argv)
{
    if(argc!=2) {
        printf("usage: ex2 nb_components\n");
        exit(0);
    }

    int size=atoi(argv[1]);
    cublasStatus_t stat;
    cublasHandle_t handle;
    stat=cublasCreate(&handle);
    int i;
    double *h_arrayA=(double*)malloc(size*sizeof(double));
    double *h_arrayB=(double*)malloc(size*sizeof(double));
    double *h_arrayC=(double*)malloc(size*sizeof(double));
    double *h_arrayCgpu=(double*)malloc(size*sizeof(double));
```

```

double *d_arrayA, *d_arrayB, *d_arrayC;

cudaMalloc((void**)&d_arrayA, size*sizeof(double));
cudaMalloc((void**)&d_arrayB, size*sizeof(double));
cudaMalloc((void**)&d_arrayC, size*sizeof(double));

for(i=0; i<size; i++) {
    h_arrayA[i]=i+1;
    h_arrayB[i]=2*(i+1);
}

unsigned int timer_cpu = 0;
cutilCheckError(cutCreateTimer(&timer_cpu));
cutilCheckError(cutStartTimer(timer_cpu));
double dot=0;
for(i=0; i<size; i++) {
    h_arrayC[i]=h_arrayA[i]+h_arrayB[i];
    dot+=(1./h_arrayC[i])*(1./h_arrayA[i]);
}
cutilCheckError(cutStopTimer(timer_cpu));
printf("CPU processing time : %f (ms) \n", cutGetTimerValue(
    timer_cpu));
cutDeleteTimer(timer_cpu);

unsigned int timer_gpu = 0;
cutilCheckError(cutCreateTimer(&timer_gpu));
cutilCheckError(cutStartTimer(timer_gpu));
stat = cublasSetVector(size, sizeof(double), h_arrayA, 1, d_arrayA,
    , 1);
stat = cublasSetVector(size, sizeof(double), h_arrayB, 1, d_arrayB,
    , 1);
int nbBlocs=(size+nbThreadsPerBloc-1)/nbThreadsPerBloc;

addition<<<nbBlocs, nbThreadsPerBloc>>>(size, d_arrayC, d_arrayA,
    d_arrayB);
inverse<<<nbBlocs, nbThreadsPerBloc>>>(size, d_arrayC);
inverse<<<nbBlocs, nbThreadsPerBloc>>>(size, d_arrayA);
double dot_gpu=0;
stat = cublasDdot(handle, size, d_arrayC, 1, d_arrayA, 1, &dot_gpu);

cutilCheckError(cutStopTimer(timer_gpu));
printf("GPU processing time : %f (ms) \n", cutGetTimerValue(
    timer_gpu));
cutDeleteTimer(timer_gpu);
printf("cpu dot %e --- gpu dot %e\n", dot, dot_gpu);

cudaFree(d_arrayA);
cudaFree(d_arrayB);
cudaFree(d_arrayC);
free(h_arrayA);
free(h_arrayB);
free(h_arrayC);
free(h_arrayCgpu);
cublasDestroy(handle);
return 0;
}

```

---

## 2.4 Third example: matrix-matrix multiplication

Matrix-matrix multiplication is an operation which is quite easy to parallelize with a GPU. If we consider that a matrix is represented using a two dimensional array,  $A[i][j]$  represents the element of the  $i^{th}$  row and of the  $j^{th}$  column. In many cases, it is easier to manipulate a 1D array instead of a 2D array. With Cuda, even if it is possible to manipulate 2D arrays, in the following we present an example based on a 1D array. For the sake of simplicity, we consider we have a square matrix of size `size`. So with a 1D array,  $A[i*size+j]$  allows us to have access to the element of the  $i^{th}$  row and of the  $j^{th}$  column.

With a sequential programming, the matrix multiplication is performed using three loops. We assume that  $A$ ,  $B$  represent two square matrices and the result of the multiplication of  $A \times B$  is  $C$ . The element  $C[i*size+j]$  is computed as follows:

$$C[i * size + j] = \sum_{k=0}^{size-1} A[i * size + k] * B[k * size + j]; \quad (2.1)$$

In Listing 2.3, the CPU computation is performed using 3 loops, one for  $i$ , one for  $j$  and one for  $k$ . In order to perform the same computation on a GPU, a naive solution consists in considering that the matrix  $C$  is split into 2 dimensional blocks. The size of each block must be chosen such as the number of threads per block is inferior to 1,024.

In Listing 2.3, we consider that a block contains 16 threads in each dimension, the variable `width` is used for that. The variable `nbTh` represents the number of threads per block. So, to be able to compute the matrix-matrix product on a GPU, each block of threads is assigned to compute the result of the product for the elements of this block. The main part of the code is quite similar to the previous code. Arrays are allocated in the CPU and the GPU. Matrices  $A$  and  $B$  are randomly initialized. Then arrays are transferred inside the GPU memory with call to `cudaMemcpy`. So the first step for each thread of a block is to compute the corresponding row and column. With a 2 dimensional decomposition, `int i= blockIdx.y*blockDim.y+ threadIdx.y;` allows us to compute the corresponding line and `int j= blockIdx.x*blockDim.x+ threadIdx.x;` the corresponding column. Then each thread has to compute the sum of the product of the line of  $A$  by the column of  $B$ . In order to use a register, the kernel `matmul` uses a variable called `sum` to compute the sum. Then the result is set into the matrix at the right place. The computation of CPU matrix-matrix multiplication is performed as described previously. A timer measures the time. In order to use 2 dimensional blocks, `dim3 dimGrid(size/width,size/width);` allows us to create `size/width` blocks in each dimension. Likewise, `dim3`

`dimBlock(width,width);` is used to create width thread in each dimension. After that, the kernel for the matrix multiplication is called. At the end of the listing, the matrix  $C$  computed by the GPU is transferred back into the CPU and we check if both matrices  $C$  computed by the CPU and the GPU are identical with a precision of  $10^{-4}$ .

With  $1,024 \times 1,024$  matrices, on a C2070M Tesla card, this code takes 37.68ms to perform the multiplication. With an Intel Xeon E31245 at 3.30GHz, it takes 2465ms without any parallelization (using only one core). Consequently the speed up between the CPU and GPU version is about 65 which is very good regarding the difficulty of parallelizing this code.

**Listing 2.3: simple Matrix-matrix multiplication with cuda**

```
#include <stdlib.h>
#include <stdio.h>
#include <string.h>
#include <math.h>
#include <assert.h>
#include "cutil_inline.h"
#include <cublas_v2.h>

const int width=16;
const int nbTh=width*width;

const int size=1024;
const int sizeMat=size*size;

__global__
void matmul(float *d_A, float *d_B, float *d_C) {
    int i= blockIdx.y*blockDim.y+ threadIdx.y;
    int j= blockIdx.x*blockDim.x+ threadIdx.x;

    float sum=0;
    for(int k=0;k<size;k++) {
        sum+=d_A[i*size+k]*d_B[k*size+j];
    }
    d_C[i*size+j]=sum;
}

int main( int argc, char** argv)
{
    float *h_arrayA=(float*)malloc(sizeMat*sizeof(float));
    float *h_arrayB=(float*)malloc(sizeMat*sizeof(float));
    float *h_arrayC=(float*)malloc(sizeMat*sizeof(float));
    float *h_arrayCgpu=(float*)malloc(sizeMat*sizeof(float));

    float *d_arrayA, *d_arrayB, *d_arrayC;

    cudaMalloc((void**)&d_arrayA,sizeMat*sizeof(float));
    cudaMalloc((void**)&d_arrayB,sizeMat*sizeof(float));
    cudaMalloc((void**)&d_arrayC,sizeMat*sizeof(float));

    srand48(32);
    for(int i=0;i<sizeMat;i++) {
        h_arrayA[i]=drand48();
    }
}
```



```

    h_arrayB[i]=drand48();
    h_arrayC[i]=0;
    h_arrayCgpu[i]=0;

}

cudaMemcpy(d_arrayA,h_arrayA, sizeMat * sizeof(float),
           cudaMemcpyHostToDevice);
cudaMemcpy(d_arrayB,h_arrayB, sizeMat * sizeof(float),
           cudaMemcpyHostToDevice);
cudaMemcpy(d_arrayC,h_arrayC, sizeMat * sizeof(float),
           cudaMemcpyHostToDevice);

unsigned int timer_cpu = 0;
cutilCheckError(cutCreateTimer(&timer_cpu));
cutilCheckError(cutStartTimer(timer_cpu));
int sum=0;
for(int i=0;i<size;i++) {
    for(int j=0;j<size;j++) {
        for(int k=0;k<size;k++) {
            h_arrayC[size*i+j]+=h_arrayA[size*i+k]*h_arrayB[size*k+j];
        }
    }
}
cutilCheckError(cutStopTimer(timer_cpu));
printf("CPU processing time : %f (ms) \n", cutGetTimerValue(
    timer_cpu));
cutDeleteTimer(timer_cpu);

unsigned int timer_gpu = 0;
cutilCheckError(cutCreateTimer(&timer_gpu));
cutilCheckError(cutStartTimer(timer_gpu));

dim3 dimGrid(size/width,size/width);
dim3 dimBlock(width,width);

matmul<<<dimGrid,dimBlock>>>(d_arrayA,d_arrayB,d_arrayC);
cudaThreadSynchronize();

cutilCheckError(cutStopTimer(timer_gpu));
printf("GPU processing time : %f (ms) \n", cutGetTimerValue(
    timer_gpu));
cutDeleteTimer(timer_gpu);

cudaMemcpy(h_arrayCgpu,d_arrayC, sizeMat * sizeof(float),
           cudaMemcpyDeviceToHost);

for(int i=0;i<sizeMat;i++)
    if (fabs(h_arrayC[i]-h_arrayCgpu[i])>1e-4)
        printf("%f %f\n",h_arrayC[i],h_arrayCgpu[i]);

cudaFree(d_arrayA);
cudaFree(d_arrayB);
cudaFree(d_arrayC);
free(h_arrayA);
free(h_arrayB);
free(h_arrayC);

```

```
    free(h_arrayCgpu);  
    return 0;  
}
```

---

---

## 2.5 Conclusion

In this chapter, three simple Cuda examples have been presented. They are quite simple. As we cannot present all the possibilities of the Cuda programming, interested readers are invited to consult Cuda programming introduction books if some issues regarding the Cuda programming are not clear.

---

## Bibliography

- [1] Jack Dongarra. Basic linear algebra subprograms technical (blast) forum standard (1). *IJHPCA*, 16(1):1–111, 2002.
- [2] J. Sanders and E. Kandrot. *CUDA by example: an introduction to general-purpose GPU programming*. Addison-Wesley, pub-AW:adr, 2010.

# Chapter 3

---

## *Ludwig: GPU implementation of a complex fluid lattice Boltzmann application*

Alan Gray and Kevin Stratford

*EPCC, The University of Edinburgh*

3.1	Introduction .....	23
3.2	Background .....	25
3.3	Single GPU Implementation .....	27
3.4	Multiple GPU Implementation .....	29
3.5	Moving Solid Particles .....	32
3.6	Summary .....	34
	Acknowledgments .....	35
	Bibliography .....	35

---

### 3.1 Introduction

The lattice Boltzmann (LB) method (for an overview see, e.g., [1]) has become a popular approach to a variety of fluid dynamics problems. It provides a way to solve the incompressible, isothermal Navier-Stokes equations and has the attractive features of being both explicit in time and local in space. This makes the LB method well-suited to parallel computation. Many efficient parallel implementations of the LB method have been undertaken, typically using a combination of distributed domain decomposition and the Message Passing Interface (MPI). However, the potential performance benefits offered by GPUs has motivated a new ‘mixed-mode’ approach to address very large problems. Here, fine-grained parallelism is implemented on the GPU, while MPI is reserved for larger-scale parallelism. This mixed mode is of increasing interest to application programmers at a time when many supercomputing services are moving toward clusters of GPU accelerated nodes. The design questions which arise when developing a lattice Boltzmann code for this type of heterogeneous system are therefore worth studying. However, similar questions also recur in many other types of stencil-based algorithms.

The first applications of LB on GPUs were to achieve fluid-like effects in computer animation, rather than scientific applications per se. These early

works include simple fluids [5], miscible two-component flow [6], and various image processing tasks based on the use of partial differential equations [7]. While these early works used relatively low level graphics APIs, the first CUDA runtime interface implementation was a two-dimensional simple fluid problem [8]. Following pioneering work on clusters of GPUs coupled via MPI to study air pollution [9], more recent work has included mixed OpenMP and CUDA [10], Posix threads and CUDA [11], and MPI and CUDA for increasingly large GPU clusters [12, 13, 14]. The heterogeneous nature of these systems has also spurred interest in approaches including automatic code generation [19] and autotuning [20] to aid application performance.

Many of these authors make use of another attractive feature of LB: the ability to include fixed solid-fluid boundary conditions as a straightforward addition to the algorithm to study, for example, flow in porous media. This points to an important application area for LB methods: that of complex fluids. Complex fluids include mixtures, surfactants, liquid crystals, and particle suspensions, and typically require additional physics beyond the bare Navier-Stokes equations to provide a full description [3]. The representation of this extra physics raises additional design questions for the application programmer. Here, we consider the *Ludwig* code [2], an LB application developed specifically for complex fluids (*Ludwig* was named for Boltzmann, 1844–1906). We will present the steps required to allow Ludwig to exploit efficiently both a single GPU, and also many GPUs in parallel. We show that Ludwig scales excellently to at least the one thousand GPU level (the largest resource available at the time of writing) with indications that the code will scale to much larger systems as they become available. In addition, we consider the steps required to represent fully-resolved moving solid particles (usually referred to as colloids in the context of complex fluids). Such particles need to have their surface resolved on the lattice scale in order to include relevant surface physics, and must be able to move: e.g., to execute Brownian motion in response to random forces from the fluid. Standard methods are available to represent such particles (e.g., [22, 23]) which are amenable to effective domain decomposition and message passing [4]. We describe below the additional considerations which arise for such moving particles when developing an implementation on a GPU.

In the following section we provide a brief overview of the lattice Boltzmann method, and mention some of the general issues which can influence performance in the CPU context. In Section 3.3, we then describe the alterations which are required to exploit the GPU architecture effectively, and highlight how the resulting code differs from the CPU version. In Section 3.4, we extend this description to include the steps required to allow exploitation of many GPUs in parallel while retaining effective scaling. We also present results for a typical benchmark for a fluid-only problem in Section 3.4 to demonstrate the success of the approach. In Section 3.5, we describe the design choices which are relevant to a GPU implementation of moving particles. Finally, we include

in the summary a number of general observations on software engineering and maintenance which arise from our experience.

### 3.2 Background

For a general complex fluid problem the starting point is the fluid velocity field  $\mathbf{u}(\mathbf{r})$ , whose evolution obeys the Navier-Stokes equations describing the conservation of mass (or density  $\rho$ ), and momentum:

$$\rho[\partial_t \mathbf{u} + (\mathbf{u} \cdot \nabla) \mathbf{u}] = -\nabla p + \eta \nabla^2 \mathbf{u} + \mathbf{f}(\mathbf{r}), \quad (3.1)$$

where  $p$  is the isotropic pressure and  $\eta$  is the viscosity. A local force  $\mathbf{f}(\mathbf{r})$  provides a means for coupling to other complex fluid constituents, e.g., it might represent the force exerted on the fluid by a curved interface between different phases or components.

The LB approach makes use of a regular three-dimensional lattice (see Figure 3.1) with discrete spacing  $\Delta r$ . It also makes use of a discrete velocity space  $\mathbf{c}_i$ , where the  $\mathbf{c}_i$  are chosen to capture the correct symmetries of the Navier-Stokes equations. A typical choice, used here, is the so-called D3Q19 basis in three dimensions where there is one velocity such that  $\mathbf{c} \Delta t$  is zero, along with six extending to the nearest neighbour lattice sites, and twelve extending to the next-nearest neighbour sites ( $\Delta t$  being the discrete time step). The fundamental object in LB is then the distribution function  $f_i(\mathbf{r}; t)$  whose moments are related to the local hydrodynamic quantities: the fluid density, momentum, and stress. The time evolution of the distribution function is described by a discrete Boltzmann equation

$$f_i(\mathbf{r} + \mathbf{c}_i \Delta t; t) - f_i(\mathbf{r}; t) = -\mathcal{L}_{ij} f_j(\mathbf{r}; t). \quad (3.2)$$

It is convenient to think of this in two stages. First, the right hand side represents the action of a collision operator  $\mathcal{L}_{ij}$ , which is local to each lattice site and relaxes the distribution toward a local equilibrium at a rate ultimately related to the fluid viscosity. Second, the left hand side represents a propagation step (sometimes referred to as streaming step), in which each element  $i$  of the distribution is displaced  $\mathbf{c}_i \Delta t$ , i.e., one lattice spacing in the appropriate direction per discrete time step.

More specifically, we store a vector of 19 double precision floating point values at each lattice site for the distribution function  $f_i(\mathbf{r}; t)$ . The collision operation  $\mathcal{L}_{ij}$ , which is local at each lattice site, may be thought of as follows. A matrix-vector multiplication  $\mathcal{M}_{ij} f_j$  is used to transform the distributions into the hydrodynamic quantities, where  $\mathcal{M}_{ij}$  is a constant 19x19 matrix related to the choice of  $\mathbf{c}_i$ . The non-conserved hydrodynamic quantities are then relaxed toward their (known) equilibrium values, and are transformed back

Listing 3.1: Collision schematic for CPU.

```

/* loop over lattice sites */
for (is = 0; is < nsites; is++) {
    ...
    /* load distribution from ``array-of-structures`` */
    for (i = 0; i < 19; i++)
        f[i] = f_aos[19*is + i]
    ...
    /* perform matrix-vector multiplication */
    for (i = 0; i < 19; i++) {
        a_tmp = 0.0;
        for (j = 0; j < 19; j++) {
            a_tmp += f[j]*M[i][j];
        }
        a[i] = a_tmp;
    }
    ...
}

```

to new post-collision distributions via the inverse transformation  $\mathcal{M}_{ij}^{-1}$ . This gives rise to the need for a minimum of  $2 \times 19^2$  floating point multiplications per lattice site. In contrast, the propagation stage consists solely of moving the distribution values between lattice sites, and involves no floating point operations.

In the CPU version, the *Ludwig* implementation stores one time level of distribution values  $f_i(\mathbf{r}; t)$ . This distribution is stored as a flat array of C data type `double`, and laid out so that all elements of the velocity are contiguous at a given site (often referred to as “array-of-structures” order). This is appropriate for sums over the distribution required locally at the collision stage as illustrated schematically in Listing 3.1: the fact that consecutive loads are from consecutive memory addresses allows the prefetcher to fully engage. (The temporary scalar `a_tmp` allows caching of the intermediate accumulated value in the innermost loop.) A variety of standard sequential optimisations are relevant for the collision stage (loop unrolling, inclusion of SIMD operations, and so on [15]). For example, the collision stage in *Ludwig* includes explicit SIMD code, which is useful if the compiler on a given platform cannot identify it. The propagation stage is separate, and is organised as a “pull” (again, various optimisation have been considered, e.g., [16, 17, 18]). No further optimisation is done here beyond ensuring that the ordering of the discrete velocities allows memory access to be as efficient as possible. While these optimisations are important, it should be remembered that for some complex fluid problems, the hydrodynamics embodied in the LB calculation is a relatively small part of the total computational cost (at the 10%-level in some cases). This means optimisation effort may be better concentrated elsewhere.

The regular 3D decomposition is illustrated in Fig. 3.1. Each local sub-




FIGURE 3.1: Left: the lattice is decomposed between MPI tasks. For clarity we show a 2D decomposition of a 3D lattice, but in practice we decompose in all 3 dimensions. Halo cells are added to each sub-domain (as shown on the upper right for a single slice) which store data retrieved from remote neighbours in the halo exchange. Lower right: the D3Q19 velocity set resident on a lattice site; highlighted are the 5 “outgoing” elements to be transferred in a specific direction.

domain is surrounded by a halo, or ghost, region of width one lattice site. While the collision is local, elements of the distribution must be exchanged at the edges of the domains to facilitate the propagation. To achieve the full 3D halo exchange, the standard approach of shifting the relevant data in each co-ordinate direction in turn is adopted. This requires appropriate synchronisation, i.e., a receive in the first co-ordinate direction must be complete before a send in the second direction involving relevant data can take place, and so on. We note that only “outgoing” elements of the distribution need to be sent at each edge. For the D3Q19 model, this reduces the volume of data traffic from 19 to 5 of the  $f_i(\mathbf{r}; t)$  per lattice site at each edge. In the CPU version, the necessary transfers are implemented in place using a vector of appropriately strided MPI datatypes for each direction.

### 3.3 Single GPU Implementation

In this section we describe the steps taken to enable *Ludwig* for the GPU. There are a number of crucial issues: first, the minimisation of data traffic between host and device; second, the optimal mapping of available parallelism onto the architecture and third, the issue of memory coalescing. We discuss each of these in turn.

While the most important section of the LB in terms of floating-point performance is the collision stage, this cannot be the only consideration for a GPU implementation. It is essential to offload all computational activity which involves the main data structures (such as the distribution) to the GPU. This includes kernels with relatively low computational demand, such as the propagation stage. All relevant data then remain resident on the GPU, to avoid expensive host-device data transfer at each iteration of the algorithm. Such transfers would negate any benefit of GPU acceleration. We note that for a complex fluid code, this requirement can extend to a considerable number of kernels, although we limit the discussion to collision and propagation for brevity.

To achieve optimal performance, it is vital to exploit fully the parallelism inherent in the GPU architecture, particularly for those matrix-vector operations within the collision kernel. The GPU architecture features a hierarchy of parallelism. At the lowest level, groups of 32 threads (warps) operate in lock-step on different data elements: this is SIMD style vector-level parallelism. Multiple warps are combined into a thread block (in which communication and synchronisation are possible), and multiple blocks can run concurrently across the streaming multiprocessors in the GPU (with no communication or synchronisation possible across blocks). To decompose on the GPU, we must choose which part of the collision to assign to which level of parallelism.

While there exists parallelism within the matrix-vector operation, the length of each vector (here, 19) is smaller than the warp size and typical thread block sizes. So we simply decompose the loop over lattice sites to all levels of parallelism, i.e., we use a separate CUDA thread for every lattice site, and each thread performs the full matrix-vector operation. We find that a block size of 256 performs well on current devices: we therefore decompose the lattice into groups of 256 sites, and assign each group to a block of CUDA threads. As the matrix  $\mathcal{M}_{ij}$  is constant, it is assigned to the fast *constant* on-chip device memory.

For the propagation stage, the GPU implementation adds a second time level of distribution values. The data-dependencies inherent in the propagation mean that the in-place propagation of the CPU version cannot be parallelised effectively without the additional time level. As both time levels may remain resident on the GPU, this is not a significant overhead.

An architectural constraint of GPUs means that optimal global memory



Listing 3.2: Collision schematic for GPU.

```

/* compute current site index 'is' from CUDA thread and block */
/* indices and load distribution from "structure-of-arrays" */

for (i = 0; i < 19; i++)
    f[i] = f_soa[nsites*i + is]

/* perform matrix-vector multiplication as before */
...

```

bandwidth is only achieved when data are structured such that threads within a *half-warp* (a group of 16 threads) load data from the same memory segment in a single transaction: this is memory coalescing. It can be seen that the array-of-structures ordering used for the distribution in the CPU code would not be suitable for coalescing; in fact, it would result in serialised memory accesses and relative poor performance. To meet the coalescing criteria and allow consecutive threads to read consecutive memory addresses on the GPU, we transpose the layout of the distribution so that, for each velocity component, consecutive sites are contiguous in memory (“structure-of-arrays” order). A schematic of the GPU collision code is given in Listing 3.2.

*Ludwig* was modified to allow a choice of distribution data layout at compilation time depending on the target architecture: CPU or GPU. We defer some further comments on software engineering aspects of the code to the summary.

### 3.4 Multiple GPU Implementation

To execute on multiple GPUs, we use the same domain decomposition and message passing framework as the CPU version. Within each sub-domain (allocated to one MPI task) the GPU implementation proceeds as described in the previous section. The only additional complication is that halo transfers between GPUs must be staged through the host (in future, direct GPU to GPU data transfers via MPI may be possible, obviating the need for these steps). This means host MPI sends must be preceded by appropriate device to host transfers and host MPI receives must be followed by corresponding host to device transfers.

In practice, this data movement requires additional GPU kernels to pack and unpack the relevant data before and after corresponding MPI calls. However, the standard shift algorithm, in which each co-ordinate direction is treated in turn, does provide some scope for the overlapping of different op-

erations. For example, after the data for the first co-ordinate direction have been retrieved by the host, these can be exchanged using MPI between hosts at the same time as kernels for packing and retrieving of data for the second co-ordinate direction are executed. This overlapping must respect the synchronisation required to ensure that data values at the corners of the sub-domain are transferred correctly. We use a separate CUDA stream for each co-ordinate direction: this allows some of the host-device communication time to be effectively “hidden” behind the host-host MPI communication, resulting in an overall speedup. The improvement is more pronounced for the smaller local lattice size, perhaps because of less CPU memory bandwidth contention. The overlapping is then a particularly valuable aid to strong scaling, where the local system size decreases as the number of GPUs increases.

To demonstrate the effectiveness of our approach, we compare the performance of both CPU and GPU versions of *Ludwig*. To test the complex fluid nature of the code, the problem is actually an immiscible fluid mixture which uses a second distribution function to introduce a composition variable. The interested reader is referred to [21] for further details. The largest total problem size used is  $2548 \times 1764 \times 1568$ . The CPU system features a Cray XE6 architecture with 2 16-core AMD Opteron CPUs per node, and with nodes interconnected using Cray Gemini technology. For GPU results, a Cray XK6 system is used: this is very similar to the XE6, but has one CPU per node replaced with an NVIDIA X2090 GPU. Each node in the GPU system therefore features a single Opteron CPU acting as a host to a single GPU. The inter-node interconnect architecture is the same as for the Cray XE6. The GPU performance tests use a prototype Cray XK6 system with 936 nodes (the largest available at the time of writing). To provide a fair comparison, we compare scaling on a *per node* basis. That is, we compare 1 fully occupied 32-core CPU node (running 32 MPI tasks) with 1 GPU node (host running 1 MPI task, and 1 device). We believe this is representative of the true “cost” of a simulation in terms of accounting budgets, or electricity.

Figure 3.2 shows the results of both weak and strong scaling tests (top and bottom panels, respectively). For weak scaling, where the local sub-domain size is kept fixed (here a relatively large  $196^3$  lattice), the time taken by an ideal code would remain constant. It can be seen that while scaling of the CPU version shows a pronounced upward slope, the GPU version scales almost perfectly. The advantage of the GPU version over the CPU version is the fact that the GPU version has a factor of 32 fewer MPI tasks per node, so communications require a significantly smaller number of larger data transfers. The performance advantage of the GPU version ranges from a factor of around 1.5 to around 1.8. Careful study by other authors [20] have found the absolute performance (in terms of floating point operations per second, or floating point operations per Watt) to be remarkably similar between architectures.

Perhaps of more interest is the strong scaling picture (lower panel in Figure 3.2) where the performance as a function of the number of nodes is measured for fixed problem size. We consider four different fixed problem sizes on both

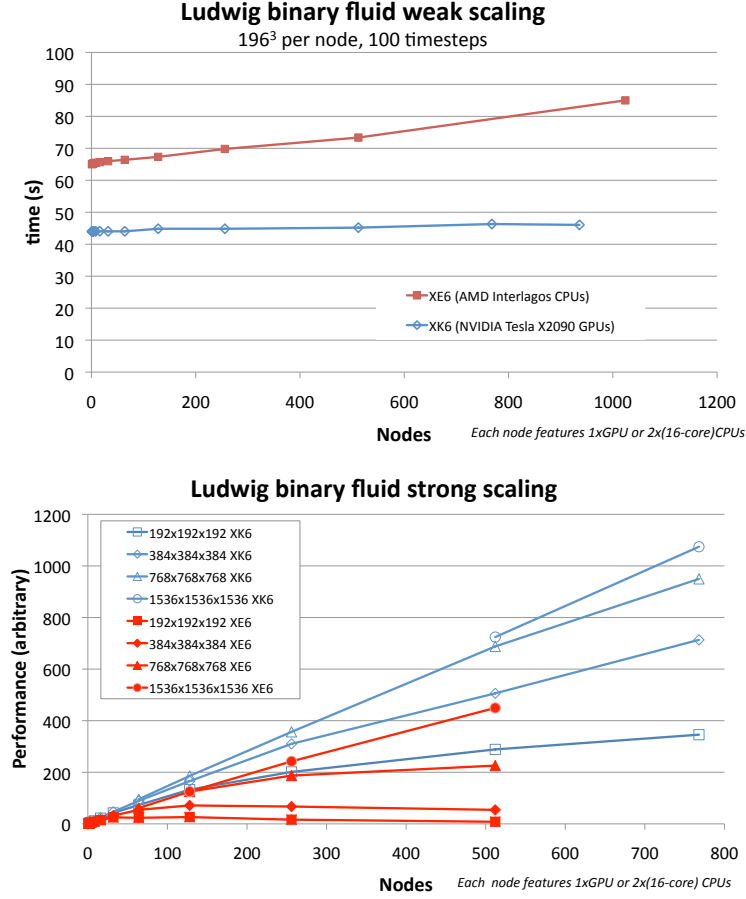


FIGURE 3.2: The weak (top) and strong (bottom) scaling of Ludwig. Closed shapes denote results using the CPU version run on the Cray XE6 (using two 16-core AMD Interlagos CPUs per node), while open shapes denote results using the GPU version on the Cray XK6 (using a single NVIDIA X2090 GPU per node). Top: the benchmark time is shown where the problem size per node is constant. Bottom: performance is shown where, for each of the four problem sizes presented, the results are scaled by the lattice volume, and all results are normalized by the same arbitrary constant.

CPU (up to 1024 nodes) and GPU (up to 768 nodes). To allow comparison, the results are scaled by total system size in each case. For strong scaling, the disparity in the number of MPI tasks is clearly revealed in the failure of the CPU version to provide any significant benefit beyond a modest number of nodes as communication overheads dominate. In contrast, the GPU version




FIGURE 3.3: A 2D illustration of colloidal particles (grey circles) interacting with the lattice, where each lattice point is represented by a cross. Left: each particle spans multiple sites. Right: the particle-facing distribution components are moved inside the particle with the direction reversed, such that the “bounce back” will be completed when the fluid propagates.

shows reasonably robust scaling even for the smaller system sizes and good scaling for the larger systems.

### 3.5 Moving Solid Particles

The introduction of moving solid particles poses an additional hurdle to efficient GPU implementation of an LB code such as *Ludwig*. In this section, we give a brief overview of the essential features of the CPU implementation, and how the considerations raised in the previous sections — maximising parallelism and minimising both host to device and GPU to GPU data movement — shape the design decisions for a GPU implementation. We restrict our discussion to a simple fluid in this section; additional solid-fluid boundary conditions (e.g., wetting at a fluid-fluid-solid contact line) usually arise elsewhere in the calculation and broadly independent of the hydrodynamic boundary conditions which are described in what follows.

Moving solid particles (here, spheres) are defined by a centre position which is allowed to move continuously across the space of the lattice, and a fixed radius which is typically a few lattice spacings  $\Delta r$ . The surface of the particle is defined by a series of *links* where a discrete velocity propagation  $\mathbf{c}_i \Delta t$  would intercept or cut the spherical shell (see Fig. 3.3). Hydrodynamic boundary conditions are then implemented via the standard approach of bounce-back on links [22, 23], where the relevant post-collision distribution values are re-

versed at the propagation stage with an appropriate correction to allow for the solid body motion. The exchange of momentum at each link must then be accumulated around the entire particle surface to provide the net hydrodynamic force and torque on the sphere. The particle motion can then be updated in a molecular dynamics-like step.

For the CPU implementation, a number of additional MPI communications are required: (1) to exchange the centre position, radius, etc of each particle as it moves and (2), to allow the accumulation of the force and torque for each particle (the links of which may be distributed between up to 8 MPI tasks in three dimensions). Appropriate marshaling of these data can provide an effective parallelisation for relatively dense particle suspensions of up to 40% solid volume fraction [4]. As a final consideration, fluid distributions must be removed and replaced consistently as a given particle moves across the lattice and changes its discrete shape.

With these features in mind, we can identify a number of competing concerns which are relevant to a GPU implementation:

1. Minimisation of host-device data transfer would argue for moving the entire particle code to the GPU. However, the code in question involves largely conditional logic (e.g., identifying cut surface links) and irregular memory accesses (e.g., access to distribution elements around a spherical particle). These operations seem poorly suited to effective parallelisation on the GPU. As an additional complication, the sums required over the particle surface would involve potentially tricky and inefficient reductions in GPU memory.
2. The alternative is to retain the relevant code on the CPU. While the transfer of the entire distribution  $f_i(\mathbf{r}; t)$  between host and device at each time step is unconscionable owing to PCIe bus bandwidth considerations, the transfer of only relevant distribution information to allow bounce-back on links is possible. At modest solid volumes only a very small fraction of the distribution data is involved in bounce-back (e.g., a 2% solid volume fraction of particles radius  $2.3\Delta r$  would involve approximately 1% of the distribution data). This option also has the advantage that no further host-device data transfers are necessary to allow the MPI exchanges required for particle information.

We have implemented the second option as follows. For each sub-domain, a list of boundary-cutting links is assembled on the CPU which includes the identity of the relevant element of the distribution. This list, together with the particle information required to compute the correct bounce-back term, are transferred to the GPU. The updates to the relevant elements of the distribution can then take place on the GPU. The corresponding information to compute the update of the particle dynamics is returned to the CPU, where the reduction over the surface links is computed. The change of particle shape may be dealt with in a similar manner: the relatively small number of

updates required at any one time step (or however frequently the particle position is updated) can be marshaled to the GPU as necessary. CAN WE SAY SOMETHING LIKE: This preliminary implementation is found to be effective on problems involving up to X per cent solid volume fraction.

Finally, we note that the CPU version actually avoids the collision calculation at solid lattice points by consulting a look-up table of solid/fluid status. On the GPU, it is perhaps preferable to perform the collision stage everywhere instead of moving the look-up table to the GPU and introducing the associated logic. Ultimately, the GPU might favour other boundary methods which treat solid and fluid on a somewhat more equal basis, for example, the immersed boundary method [24, 25] or smoothed profile method [26]. However, the approach adopted here allows us to exploit the GPU for the intensive fluid simulation whilst maintaining the complex code required for particles on the CPU. Overheads of CPU-GPU transfer are minimised by transferring only those data relevant to the hydrodynamic interaction implemented via bounce-back on links.

It is perhaps interesting at this point to make some more general observations on the software engineering challenge presented when extending an existing CPU code to the GPU. The already complex task of maintaining the code in a portable fashion while also maintaining performance is currently formidable. To help this process, we have followed a number of basic principles. First, in order to port to the GPU in an incremental fashion, we have tried to maintain the modular structure of the CPU where possible. For each data structure, such as the distribution, a separate analogue is maintained in both the CPU and GPU memory spaces. However, the GPU copy does not include the complete CPU structure: in particular, non-intrinsic datatypes such as MPI datatypes are not required on the GPU. Functions to marshal data between CPU and GPU are provided for each data structure, abstracting the underlying CUDA implementation. (This reasonably lightweight abstraction layer could also allow an OpenCL implementation to be developed.) This makes it easy to switch between the CPU and GPU for different components in the code, which is useful in development and testing. GPU functionality can be added incrementally while retaining a code that runs correctly (albeit slowly due to data transfer overheads). Once all relevant components are moved to the GPU, it becomes possible to remove such data transfers and keep the entire problem resident on the device.

---

### 3.6 Summary

The *Ludwig* LB application, which specializes in complex fluid problems, has undergone the developments required to use large-scale parallel GPU accelerated supercomputers effectively. The new functionality augments the orig-

inal code, so that *Ludwig* can run on either CPU-based or GPU-accelerated systems. Using the NVIDIA CUDA programming language, all the computational kernels were offloaded to the GPU. This included not only those responsible for the majority of the compute time, but any that accessed the relevant data structures. This allows data to be kept resident on the GPU and avoids expensive data transfers. We described steps taken to optimise the performance on the GPU architecture by reducing off-chip memory accesses and restructuring the data layout in memory to ensure the available memory bandwidth was exploited fully.

A new halo-exchange communication phase for the code was developed to allow efficient parallel scaling to many GPUs. The CPU version relies on MPI datatype functionality for CPU to CPU communication; we have explicitly written buffer packing/unpacking and transfer operations to for GPU to GPU data exchange, staged through the host CPUs. This includes the filtering of velocity components to only transfer those required in a specific direction, hence reducing the overall volume of data transferred. We used CUDA stream functionality to overlap separate stages within the communication phase, and reduce the overall communication time. The effect was measured to be more profound for smaller local system sizes, i.e. it is especially useful in aiding strong scaling.

We described work to enable colloidal particles in the simulation with minimal overhead. We implemented these in such a way that we avoid a major diversion of the CPU and GPU codebases, whilst minimising data transfers between the CPU and GPU at each timestep. We keep the majority of the (relatively inexpensive) particle related code on the CPU, while offloading only those parts responsible for the interaction with the fluid to the GPU (where the fluid data resides).

Our resulting package is able to exploit the largest of GPU accelerated architectures to address a wide range of complex problems.

---

## Acknowledgments

AG was supported by the CRESTA project, which has received funding from the European Community's Seventh Framework Programme (ICT-2011.9.13) under Grant Agreement no. 287703. KS was supported by UK EPSRC under grant EV/J007404/1.

---

## Bibliography

- [1] S. Succi, *The lattice Boltzmann equation and beyond*, Oxford University Press, Oxford, 2001.
- [2] Desplat, J.-C., I. Pagonabarraga, and P. Bladon, *LUDWIG: A parallel lattice-Boltzmann code for complex fluids*. Comput. Phys. Comms., **134**, 273, 2001.
- [3] C.K. Aidun and J.R. Clausen, *Lattice Boltzmann method for complex flows*, Ann. Rev. Fluid Mech., **42** 439–472 (2010).
- [4] K. Stratford and I. Pagonabarraga, Parallel domain decomposition for lattice Boltzmann with moving particles, CAMWA REF **55**, 1585 (2008).
- [5] X. Wei, W. Li, K. Müller, and A.E. Kaufman, *The lattice Boltzmann method for simulating gaseous phenomena*, IEEE Transactions on Visualization and Computer Graphics, **10**, 164–176 (2004).
- [6] H. Zhu, X. Liu, Y. Liu, and E. Wu, *Simulation of miscible binary mixtures based on lattice Boltzmann method*, Comp. Anim. Virtual Worlds, **17**, 403–410 (2006).
- [7] Y. Zhao, *Lattice Boltzmann based PDE solver on the GPU*, Visual Comput., doi 10.1007/s00371-0070191-y (2007).
- [8] J. Tölke, Implementation of a lattice Boltzmann kernel using the compute unified device architecture developed by nVIDIA, Comput. Visual Sci. **13** 29–39 (2010).
- [9] Z. Fan, F. Qiu, A. Kaufman, and S. Yoakum-Stover, *GPU cluster for high performance computing*, Proceedings of ACM/IEEE Supercomputing Conference, pp. 47–59, IEEE Computer Society Press, Pittsburgh, PA (2004).
- [10] J. Myre, S.D.C. Walsh, D. Lilja, and M.O. Saar, *Performance analysis of single-phase, multiphase, and multicomponent lattice Boltzmann fluid flow simulations on GPU clusters*, Concurrency Computat.: Pract. Exper., **23**, 332–350 (2011).
- [11] C. Obrecht, F. Kuznik, B. Tourancheau, and J.-J. Roux, *Multi-GPU implementation of the lattice Boltzmann method*, Comput. Math. with Applications, doi:10.1016/j.camwa.2011.02.020 (2011).
- [12] M. Bernaschi, M. Fatica, S. Melchionna, S. Succi, and E. Kaxiras, *A flexible high-performance lattice Boltzmann GPU code for the simulations of fluid flow in complex geometries*, Concurrency Computat.: Pract. Exper., **22**, 1–14 (2010).



- [13] W. Xian and A. Takayuki, *Multi-GPU performance of incompressible flow computation by lattice Boltzmann method on GPU cluster*, Parallel Comput., doi:10.1016/j.parco.2011.02.007 (2011).
- [14] C. Feichtinger, J. Habich, H. Köstler, G. Hager, U. Rüde, and G. Wellein, A flexible patch-based lattice Boltzmann parallelization approach for heterogeneous GPU-CPU clusters, *Parallel Computing* **37** 536–549 (2011).
- [15] G. Wellein, T. Zeiser, G. Hager, and S. Donath, On the single processor performance of simple lattice Boltzmann kernels, *Computers and Fluids*, **35**, 910–919 (2006).
- [16] T. Pohl, M. Kowarschik, J. Wilke, K. Igelberger, and U. Rüde, Optimization and profiling of the cache performance of parallel lattice Boltzmann code, *Parallel Process Lett.* **13** 549–560 (2003).
- [17] K. Mattila, J. Hyväluoma, T. Rossi, M. Aspö and J. Westerholm, An efficient swap algorithm for the lattice Boltzmann method, *Comput. Phys. Comms.* **176** 200–210 (2007).
- [18] M. Wittmann, T. Zeiser, G. Hager, and G. Wellein, Comparison of different propagation steps for lattice Boltzmann methods, *Comput. Math with Appl.* doi:10.1016/j.camwa.2012.05.002 (2012).
- [19] S.D.C. Walsh and M.O. Saar, Developing extensible lattice Boltzmann simulators for general-purpose graphics-processing units, *Comm. Comput. Phys.*, **13** 867–879 (2013).
- [20] S. Williams, L. Oliker, J. Carter, and J. Shalf, Extracting ultra-scale lattice Boltzmann performance via hierarchical and distributed auto-tuning, *Proc. SC2011*.
- [21] K. Stratford, R. Adhikari, I. Pagonabarraga, and J.-C. Desplat, *Lattice Boltzmann for Binary Fluids with Suspended Colloids*, J. Stat. Phys. **121**, 163 (2005).
- [22] A.J.C. Ladd, Numerical simulations of particle suspensions via a discretized Boltzmann equation. Part 1. Theoretical foundation, *J. Fluid Mech.* **271** 285–309 (1994); Part II. Numerical results, *ibid.* **271** 311–339 (1994).
- [23] N.-Q. Nguyen and A.J.C. Ladd, Lubrication corrections for lattice Boltzmann simulations of particle suspensions, *Phys. Rev. E* **66** 046708 (2002).
- [24] C.S. Peskin, Flow patterns around heart valves; a numerical method, *J. Comp. Phys.*, **10**, 252–271 (1972); C.S. Peskin, The immersed boundary method, *Acta Numerica* **11** 479–517 (2002).

- [25] Z.-G. Feng and E.E. Michaelides, The immersed boundary-lattice Boltzmann method for solving fluid-particles interaction problem, *J. Comp. Phys.*, **195** 602–628 (2004).
- [26] Y. Nakayama and R. Yamamoto, Simulation method to resolve hydrodynamic interactions in colloidal dispersions, *Phys. Rev. E*, **71** 036707 (2005).