

1

# targetDP Specification

2

Alan Gray and Kevin Stratford

# Contents

1	<b>1 Introduction</b>	<b>2</b>
2	1.1 Abstract . . . . .	2
3	1.2 Motivation . . . . .	2
4	1.3 Glossary . . . . .	2
5	<b>2 Execution model</b>	<b>3</b>
6	<b>3 Memory model</b>	<b>5</b>
7	3.1 Model overview . . . . .	5
8	3.2 Implementation . . . . .	6
9	<b>4 Memory Management</b>	<b>7</b>
10	4.1 targetMalloc . . . . .	8
11	4.2 targetCalloc . . . . .	9
12	4.3 targetFree . . . . .	10
1	4.4 copyToTarget . . . . .	11
2	4.5 copyFromTarget . . . . .	12
3	4.6 targetConst . . . . .	13
4	4.7 copyConstToTarget . . . . .	14
5	4.8 copyConstFromTarget . . . . .	15
6	4.9 targetConstAddress . . . . .	16
7	4.10 targetInit3D . . . . .	17
8	4.11 targetFinalize3D . . . . .	18
9	4.12 copyToTargetBoundary3D . . . . .	19
10	4.13 copyFromTargetBoundary3D . . . . .	20

---

## CONTENTS

---

11	4.14 copyToTargetPointerMap3D . . . . .	21
12	4.15 copyFromTargetPointerMap3D . . . . .	22
13	<b>5 Data Parallel Execution</b>	<b>23</b>
14	5.1 targetEntry . . . . .	24
15	5.2 target . . . . .	25
16	5.3 targetHost . . . . .	26
17	5.4 targetLaunch . . . . .	27
18	5.5 targetSynchronize . . . . .	28
19	5.6 targetTLP . . . . .	29
20	5.7 targetILP . . . . .	30
21	5.8 targetCoords3D . . . . .	31
22	5.9 targetIndex3D . . . . .	32
23	<b>6 Examples</b>	<b>33</b>

# Chapter 1

## Introduction

COMMENT: TO DO.

### 1.1 Abstract

\*\*

### 1.2 Motivation

\*\*

### 1.3 Glossary

- Accelerator: \*\*\*

- Host: \*\*\*

- Target: \*\*\*

- CUDA: \*\*\*

## 9 Chapter 2

# 10 Execution model

11 The terminology “host” is used to refer to the CPU that is hosting the execution  
12 of the application, and “target” to refer to the device targeted for execution of  
13 data parallel operations. The target can be the same CPU as the host, or it can  
14 be a separate device such as an accelerator.

15 The targetDP API follows the fork-join model of parallel execution. When  
16 the application initiates, a single thread executes sequentially on the host, until  
17 it encounters a function to be launched on the target. This function will be exe-  
18 cuted by a team of threads on the target cooperating in a data-parallel manner  
19 (e.g. for a structured grid problem each thread is responsible for a subset of the  
20 grid).

21 To ensure that the target function has completed, a `targetSynchronize`  
1 statement should follow the code location where target function is launched.  
2 When the initial thread encounters this statement, it will wait until the target  
3 region has completed. It is possible for the initial thread to execute other in-  
4 structions (which do not depend on the results of the target function), after the  
5 function launch but before the synchronisation call. This may result in overlap-  
6 ping of host and target execution, and hence optimisation, in some implemen-  
7 tations. Once the target function has completed, the initial thread will continue  
8 sequentially until another target function launch is encountered.

9 Within each target function, each thread is given a unique index which it  
10 uses to work in a data-parallel manner. Each thread works independently from  
11 all others, but usually operating on a shared data structure where the index is  
12 used to determine the portion of data to process.

13 **COMMENT: Possibly add `targetSyncThreads()` to allow communication within**

---

14 target functions?]

## 15 Chapter 3

# 16 Memory model

### 17 3.1 Model overview

18 The targetDP model draws a distinction between the memory spaces accessed  
19 by the host and target, such that code executed on the host always accesses the  
20 host memory space, and code executed on the target (i.e. within target func-  
21 tions) always accesses the target memory space. The host memory space can be  
22 initialised using regular C/C++ functionality, and the targetDP API provides the  
1 functionality necessary to manage the target memory space and transfer data  
2 between the host and target. For each data-parallel data structure, the program-  
1 mer should create both host and target copies, and should update these from  
2 each other as and when required.

#### 3 3.1.1 Host memory model

4 The sequential thread executing host code should always access host data struc-  
5 tures. The memory model is the same as one would expect from a regular  
6 sequential application.

#### 7 3.1.2 Target memory model

8 The team of threads performing the execution of target functions should always  
9 act on target data structures. These can take one of 3 forms:

- 10 1. Those created using the targetDP memory allocation API functions. These  
11 are shared between all threads in the team, where each thread should use

- 12        its unique index to access the portion of data for which it is responsible.
- 13        2. Those created using the targetDP constant data management functionality.
- 14        These are read-only and normally used for relatively small amounts of
- 1        constant data.
- 2        3. Those declared within the body of a target function. These are private to
- 3        each thread in the team and should be used for temporary scratch struc-
- 4        tures.
- 5        **COMMENT:** For 3, at the moment any data declared in the function but
- 6        above the targetTLP keyword will be private for GPU threads but shared for
- 7        CPU (since outside parallel region). So either need to make this clearer above,
- 8        or somehow move OpenMP parallel region to target launch.

## 9        3.2 Implementation

### 10       3.2.1 C

11       The target memory structures will exist on the same physical memory as the host

12       structures. The implementation may either use separate target copies (managed

13       using regular C/C++ memory management functionality), or use pointer alias-

14       ing for the target versions such that a reference to any part of a target structure

15       will correspond to exactly the same physical address as that of the correspond-

1       ing host structure.

### 2       3.2.2 CUDA

3       The target memory space will exist on the distinct GPU memory, i.e. in a sepa-

4       rate memory space from the host structures.



## <sup>5</sup> **Chapter 4**

# <sup>6</sup> **Memory Management**

## 7 **4.1 targetMalloc**

### 8 **4.1.1 Description**

9 The targetMalloc function allocates memory on the target.

### 10 **4.1.2 Syntax**

11 `void targetMalloc(void **targetPtr, size_t n);`

- 12     • targetPtr: A pointer to the allocated memory.
- 13     • n: The number of bytes to be allocated.

### 1 **4.1.3 Example**

2 See Line 1 in Figure 6.2 in Section 6.

### 3 **4.1.4 Implementation**

#### 4 **C**

5 `malloc`

#### 6 **CUDA**

7 `cudaMalloc`

## 8 **4.2 targetCalloc**

### 9 **4.2.1 Description**

10 The `targetCalloc` function allocates, and initializes to zero, memory on the  
11 target.

### 12 **4.2.2 Syntax**

13 `void targetCalloc(void **targetPtr, size_t n);`

- 14 • `targetptr`: A pointer to the allocated memory.
- 15 • `n`: The number of bytes to be allocated.

### 1 **4.2.3 Example**

2 Analogous to Line 1 in Figure 6.2 in Section 6.

### 3 **4.2.4 Implementation**

#### 4 **C**

5 `calloc`

#### 6 **CUDA**

7 `cudaMalloc` followed by `cudaMemset`

## 8 **4.3 targetFree**

### 9 **4.3.1 Description**

10 The targetFree function deallocates memory on the target.

### 11 **4.3.2 Syntax**

12 void targetFree(void \*targetPtr);

- 13     • targetPtr: A pointer to the memory to be freed.

### 14 **4.3.3 Example**

15 See Line 10 in Figure 6.2 in Section 6.

### 1 **4.3.4 Implementation**

2 **C**

3 free

4 **CUDA**

5 cudaFree

## 6 4.4 copyToTarget

### 7 4.4.1 Description

8 The copyToTarget function copies data from the host to the target.

### 9 4.4.2 Syntax

10 `void copyToTarget(void *targetData, const void *data, size_t n);`

- 11     • targetData: A pointer to the destination array on the target.
- 12     • data: A pointer to the source array on the host.
- 13     • n: The number of bytes to be copied.

### 14 4.4.3 Example

15 See Line 3 in Figure 6.2 in Section 6.

### 16 4.4.4 Implementation

#### 1 C

2 `memcpy`

#### 3 CUDA

4 `cudaMemcpy`

## 5 4.5 copyFromTarget

### 6 4.5.1 Description

7 The copyFromTarget function copies data from the target to the host.

### 8 4.5.2 Syntax

9 `void copyFromTarget(void *data, const void *targetData, size_t n);`

- 10 • data: A pointer to the destination array on the host.
- 11 • targetData: A pointer to the source array on the target.
- 12 • n: The number of bytes to be copied.

### 13 4.5.3 Example

14 See Line 9 in Figure 6.2 in Section 6.

### 15 4.5.4 Implementation

#### 16 C

17 `memcpy`

#### 18 CUDA

1 `cudaMemcpy`

## 2   **4.6   targetConst**

### 3   **4.6.1   Description**

4   The `__targetConst__` keyword is used in a variable or array declaration to  
5   specify that the corresponding data can be treated as constant (read-only) on  
6   the target.

### 7   **4.6.2   Syntax**

8   `__targetConst__ type variableName`

- 9       • `variableName`: The name of the variable or array.
- 10      • `type`: The type of variabe or array.

### 11   **4.6.3   Example**

12   `**TO DO**`.

### 13   **4.6.4   Implementation**

14   **C**

15   Holds no value

16   **CUDA**

17   `__constant__`

## 1 **4.7 copyConstToTarget**

### 2 **4.7.1 Description**

3 The `copyConstToTarget` function copies data from the host to the target, where  
4 the data will remain constant (read-only) during the execution of functions on  
5 the target.

### 6 **4.7.2 Syntax**

7 `void copyConstToTarget(void *targetData, const void *data, size_t n);`

- 8 • `targetData`: A pointer to the destination array on the target. This must  
9 have been declared using the `__targetConst__` keyword.
- 10 • `data`: A pointer to the source array on the host.
- 11 • `n`: The number of bytes to be copied.

### 12 **4.7.3 Example**

13 See Line 4 in Figure 6.2 in Section 6.

### 14 **4.7.4 Implementation**

#### 15 **C**

16 `memcpy`

#### 17 **CUDA**

1 `cudaMemcpyToSymbol`



## 2 **4.8 copyConstFromTarget**

### 3 **4.8.1 Description**

4 The `copyConstFromTarget` function copies data from a constant data location  
5 on the target to the host.

### 6 **4.8.2 Syntax**

7 `void copyConstToTarget(void *targetData, const void *data, size_t n);`

- 8     • `data`: A pointer to the destination array on the host.
- 9     • `targetData`: A pointer to the source array on the target. This must have  
10       been declared using the `__targetConst__` keyword.
- 11     • `n`: The number of bytes to be copied.

### 12 **4.8.3 Example**

13 Analogous to Line 4 in Figure 6.2 in Section 6.

### 14 **4.8.4 Implementation**

#### 15 **C**

16 `memcpy`

#### 17 **CUDA**

1 `cudaMemcpyFromSymbol`

## 2 **4.9 targetConstAddress**

### 3 **4.9.1 Description**

4 The `targetConstAddress` function provides the target address for a constant  
5 object.

### 6 **4.9.2 Syntax**

7 `void targetConstAddress(void **address, objectType object);`

- 8     • `address` (output): The pointer to the constant object on the target.
- 9     • `objectType`: The type of the object.
- 10    • `object` (input): The constant object on the target. This should have been  
11      declared using the `__targetConst__` keyword.

### 12 **4.9.3 Example**

1 `**TO DO**.`

### 2 **4.9.4 Implementation**

#### 3 **C**

4 Explicit copying of address.

#### 5 **CUDA**

6 `cudaGetSymbolAddress`

## 7 **4.10 targetInit3D**

### 8 **4.10.1 Description**

9 The targetInit3D initialises the environment required to perform any of fol-  
10 lowing the targetDP 3D lattice operations.

### 11 **4.10.2 Syntax**

12 void targetInit3D(size\_t extent, size\_t nFields);

- 13 • extent: The total extent of data parallelism (e.g. the number of lattice  
14 sites).
- 15 • nFields: The extent of data resident within each parallel partition (e.g.  
16 the number of fields per lattice site).

### 17 **4.10.3 Example**

18 **\*\*TO DO\*\*.**

### 19 **4.10.4 Implementation**

#### 20 **C**

21 **\*\*TO DO\*\*.**

#### 1 **CUDA**

2 **\*\*TO DO\*\*.**

## **3 4.11 targetFinalize3D**

### **4 4.11.1 Description**

5 The targetFinalize3D finalizes the targetDP 3D environment.

### **6 4.11.2 Syntax**

7 `void targetFinalize3D();`

### **8 4.11.3 Example**

9 `**TO DO**.`

### **10 4.11.4 Implementation**

11 **C**

12 `**TO DO**.`

13 **CUDA**

14 `**TO DO**.`

## 15 **4.12 copyToTargetBoundary3D**

### 16 **4.12.1 Description**

17 The `copyToTargetBoundary3D` function copies the data corresponding to the  
18 boundaries of a 3D lattice from the host to the target.

### 19 **4.12.2 Syntax**

20 `void copyToTargetBoundary3D(void *targetData, const void *data, size_t extent3D[3], size_t nFields, size_t offset, size_t depth)`

- 21     • `targetData`: A pointer to the destination array on the target.
- 1     • `data`: A pointer to the source array on the host.
- 2     • `extent3D`: An array of 3 integers corresponding to the 3D dimensions of  
3         the lattice.
- 4     • `nFields`: The number of fields per lattice site.
- 5     • `offset`: The number of sites from the lattice edge at which each boundary  
6         face should start.
- 7     • `depth`: The depth of each boundary face.

### 8 **4.12.3 Example**

9 `**TO DO**`.

### 10 **4.12.4 Implementation**

11 **C**

12 `**TO DO**`.

13 **CUDA**

14 `**TO DO**`.

## 15 **4.13 copyFromTargetBoundary3D**

### 16 **4.13.1 Description**

17 The copyFromTargetBoundary3D function copies the data corresponding to the  
18 boundaries of a 3D lattice from the target to the host.

### 19 **4.13.2 Syntax**

20 `void copyFromTargetBoundary3D(void *data, const void *targetData, size_t extent3D[3], size_t`

- 21     • data: A pointer to the destination array on the host.
- 22     • targetData: A pointer to the source array on the target.
- 23     • extent3D: An array of 3 integers corresponding to the 3D dimensions of  
24         the lattice.
- 25     • nFields: The number of fields per lattice site.
- 26     • offset: The number of sites from the lattice edge at which each boundary  
1         face should start.
- 2     • depth: The depth of each boundary face.

### 3 **4.13.3 Example**

4 `**TO DO**.`

### 5 **4.13.4 Implementation**

6 **C**

7 `**TO DO**.`

8 **CUDA**

9 `**TO DO**.`

## 10 **4.14 copyToTargetPointerMap3D**

### 11 **4.14.1 Description**

12 The `copyToTargetPointerMap3D` function copies a subset of lattice data from  
13 the host to the target. The sites to be included are defined using an array of  
14 pointers passed as input.

### 15 **4.14.2 Syntax**

```
16 void copyToTargetPointerMap3D(void *targetData, const void *data,  
17                               size_t extent3D[3], size_t nField,  
18                               int includeNeighbours, void** pointerArray);
```

- 19 • `targetData`: A pointer to the destination array on the target.
- 20 • `data`: A pointer to the source array on the host.
- 21 • `extent3D`: An array of 3 integers corresponding to the 3D dimensions of  
22 the lattice.
- 23 • `nField`: The number of fields per lattice site.
- 24 • `includeNeighbours`: A boolean switch to specify whether each included  
25 site should also have its neighbours included (in the 19-point 3D stencil).
- 26 • `pointerArray`: An array of `nSite` pointers, where `nSite` is the total num-  
1 ber of lattice sites. Each lattice site should be included unless the pointer  
2 corresponding to that site is `NULL`.

### 1 **4.14.3 Example**

2 **\*\*TO DO\*\***.

### 3 **4.14.4 Implementation**

#### 4 **C**

5 **\*\*TO DO\*\***.

#### 6 **CUDA**

7 **\*\*TO DO\*\***.

## 8 **4.15 copyFromTargetPointerMap3D**

### 9 **4.15.1 Description**

10 The `copyFromTargetPointerMap3D` function copies a subset of lattice data from  
11 the target to the host. The sites to be included are defined using an array of  
12 pointers passed as input.

### 13 **4.15.2 Syntax**

```
14 void copyFromTargetPointerMap3D(void *data, const void *targetData,  
15                                size_t extent3D[3], size_t nField,  
16                                int includeNeighbours, void** pointerArray);
```

- 1 • `data`: A pointer to the destination array on the host.
- 2 • `targetData`: A pointer to the source array on the target.
- 3 • `extent3D`: An array of 3 integers corresponding to the 3D dimensions of  
4 the lattice.
- 5 • `nField`: The number of fields per lattice site.
- 6 • `includeNeighbours`: A boolean switch to specify whether each included  
7 site should also have its neighbours included (in the 19-point 3D stencil).
- 8 • `pointerArray`: An array of `nSite` pointers, where `nSite` is the total num-  
9 ber of lattice sites. Each lattice site should be included unless the pointer  
10 corresponding to that site is `NULL`.

### 11 **4.15.3 Example**

12 **\*\*TO DO\*\***.

### 13 **4.15.4 Implementation**

#### 14 **C**

15 **\*\*TO DO\*\***.

#### 16 **CUDA**

1 **\*\*TO DO\*\***.



## <sup>2</sup> **Chapter 5**

# <sup>3</sup> **Data Parallel Execution**

## 4 **5.1 targetEntry**

### 5 **5.1.1 Description**

6 The `__targetEntry__` keyword is used in a function declaration or definition  
7 to specify that the function should be compiled for the target, and that it will be  
8 called directly from host code.

### 9 **5.1.2 Syntax**

10 `__targetEntry__` `functionReturnType` `functionName`(...

- 11 • `functionName`: The name of the function to be compiled for the target.
- 12 • `functionReturnType`: The return type of the function.
- 13 • ... the remainder of the function declaration or definition.

### 14 **5.1.3 Example**

15 See Line 1 in Figure 6.1 in Section 6.

### 16 **5.1.4 Implementation**

1 **C**

2 Holds no value.

3 **CUDA**

## 4 **5.2 target**

### 5 **5.2.1 Description**

6 The `__target__` keyword is used in a function declaration or definition to spec-  
7 ify that the function should be compiled for the target, and that it will be called  
8 from a `targetEntry` or another `target` function.

### 9 **5.2.2 Syntax**

10 `__target__` `functionReturnType` `functionName`(...

- 11 • `functionName`: The name of the function to be compiled for the target.
- 12 • `functionReturnType`: The return type of the function.
- 13 • ... the remainder of the function declaration or definition.

### 14 **5.2.3 Example**

15 Analogous to Line 1 in Figure 6.1 in Section 6.

### 16 **5.2.4 Implementation**

17 **C**

18 Holds no value.

1 **CUDA**

## 2 **5.3 targetHost**

### 3 **5.3.1 Description**

4 The `__targetHost__` keyword is used in a function declaration or definition to  
5 specify that the function should be compiled for the host.

### 6 **5.3.2 Syntax**

7 `__targetHost__` functionReturnType functionName(...

- 8     • `functionName`: The name of the function to be compiled for the host.
- 9     • `functionReturnType`: The return type of the function.
- 10    • ... the remainder of the function declaration or definition.

### 11 **5.3.3 Example**

12 Analogous to Line 1 in Figure 6.1 in Section 6.

### 13 **5.3.4 Implementation**

#### 1 **C**

2 Holds no value.

#### 3 **CUDA**

4 `extern 'C' __host__`

## 5.4 targetLaunch

### 5.4.1 Description

The `__targetLaunch__` syntax is used to launch a function across a data parallel target architecture.

### 5.4.2 Syntax

```
functionName __targetLaunch__(size_t extent) \  
    (functionArgument1,functionArgument2,...);
```

- `functionName`: The name of the function to be launched. This function must be declared as `__targetEntry__`.
- `functionArguments`: The arguments to the function `functionName`
- `extent`: The total extent of data parallelism.

### 5.4.3 Example

See Line 6 in Figure 6.2 in Section 6.

### 5.4.4 Implementation

#### C

Holds no value.

#### CUDA

CUDA `<<<...>>>` syntax.

## 23 **5.5 targetSynchronize**

### 1 **5.5.1 Description**

2 The targetSynchronize function is used to block until the preceeding \_\_targetLaunch\_\_  
3 has completed.

### 4 **5.5.2 Syntax**

5 `void targetSynchronize();`

### 6 **5.5.3 Example**

7 See Line 7 in Figure 6.2 in Section 6.

### 8 **5.5.4 Implementation**

#### 9 **C**

10 Dummy function.

#### 11 **CUDA**

12 `cudaThreadSynchronize`

## 13 5.6 targetTLP

### 14 5.6.1 Description

15 The `__targetTLP__` syntax is used, within a `__targetEntry__` function, to  
16 specify that the proceeding block of code should be executed in parallel and  
17 mapped to thread level parallelism (TLP). Note that the behaviour of this op-  
18 eration depends on the defined virtual vector length (VVL), which controls the  
19 lower-level Instruction Level Parallelism (ILP) (see following section).

### 20 5.6.2 Syntax

```
21 __targetTLP__(int baseIndex, size_t extent)  
22 {  
  1 //code to be executed in parallel  
  2 }
```

- 3     • extent: The total extent of data parallelism, including both TLP and ILP
- 4     • baseIndex: the TLP index. This will vary from 0 to extent-VVL with  
5       stride VVL. This index should be combined with the ILP index to access  
6       shared arrays within the code block (see following section).

### 7 5.6.3 Example

8 See Line 4 in Figure 6.1 in Section 6.

### 9 5.6.4 Implementation

#### 10 C

11 OpenMP parallel loop.

#### 12 CUDA

13 CUDA thread lookup.

## 14 5.7 targetILP

### 15 5.7.1 Description

16 The `__targetILP__` syntax is used, within a `__targetTLP__` region, to specify  
17 that the proceeding block of code should be executed in parallel and mapped to  
18 instruction level parallelism (ILP), where the extent of the ILP is defined by the  
1 virtual vector length (VVL) in the targetDP implementation. **COMMENT: need**  
2 **to document VVL in more detail somewhere and link.**

### 3 5.7.2 Syntax

```
4 __targetILP__(int vecIndex)
5 {
6     //code to be executed in parallel
7 }
```

- 8 • `baseIndex`: the ILP index. This will vary from 0 to VVL-1. This index  
9 should be combined with the TLP index to access shared arrays within the  
10 code block (see previous section).

### 11 5.7.3 Example

12 See Line 9 in Figure 6.1 in Section 6.

### 13 5.7.4 Implementation

#### 14 C

15 Short vectorizable loop.

#### 16 CUDA

17 Short vectorizable loop.



## 18 **5.8 targetCoords3D**

### 1 **5.8.1 Description**

2 The targetCoords3D function provides the 3D lattice coordinates corresponding  
3 to a specified linear index.

### 516 **5.8.2 Syntax**

517 `void targetCoords3D(int coords3D[3], int extent3D[3], int index);`

- 518     • `coords3D` (output): an array of 3 integers to be populated with ther 3D  
519       coordinates.
- 520     • `extent3D` (input): An array of 3 integers corresponding to the 3D dimen-  
521       sions of the lattice.
- 522     • `index` (input): the linear index.

### 523 **5.8.3 Example**

524 `**TO DO**`

### 525 **5.8.4 Implementation**

#### 526 **C**

527 `**TO DO**`

#### 528 **CUDA**

529 `**TO DO**`

## 530 **5.9 targetIndex3D**

### 531 **5.9.1 Description**

532 The targetIndex3D function returns the linear index corresponding to a speci-  
533 fied set of 3D lattice coordinates.

### 534 **5.9.2 Syntax**

535 `int targetIndex3D(int Xcoord,int Ycoord,int Zcoord,int extent3D[3]);`

- 536 • Xcoord (input): the specified coordinate in the X direction.
- 537 • Ycoord (input): the specified coordinate in the Y direction.
- 538 • Zcoord (input): the specified coordinate in the Z direction.
- 539 • extent3D (input): an array of 3 integers corresponding to the 3D dimen-  
540 sions of the lattice.

### 541 **5.9.3 Example**

542 `**TO DO**`

### 543 **5.9.4 Implementation**

544 **C**

545 `**TO DO**`

546 **CUDA**

547 `**TO DO**`

## Chapter 6

## Examples

COMMENT: TO DO: describe this example.

```
__targetEntry__ void scale(double* t_field) {  
    int baseIndex;  
    __targetTLP__(baseIndex, N) {  
        int iDim, vecIndex;  
        for (iDim = 0; iDim < 3; iDim++) {  
            __targetILP__(vecIndex) \  
                t_field[iDim*N + baseIndex + vecIndex] = \  
                t_a*t_field[iDim*N + baseIndex + vecIndex];  
        }  
    }  
    return;  
}
```

Figure 6.1: \*\*\*

---

```
targetMalloc((void **) &t_field , datasize); 1
copyToTarget(t_field , field , datasize);      2
copyConstToTarget(&t_a , &a, sizeof(double)); 3
scale __targetLaunch__(N) (t_field);           4
targetSynchronize();                           5
scale __targetLaunch__(N) (t_field);           6
targetSynchronize();                           7
copyFromTarget(field , t_field , datasize);    8
targetFree(t_field);                          9
targetFree(t_field);                         10
```

Figure 6.2: \*\*\*